



Lecture 2

Linear Algebra and Probability

Review

CSE465: Pattern Recognition and Neural Network

Sec: 3

Faculty: Silvia Ahmed (SvA)

Summer 2025

Today's Topics

- Linear Algebra Essentials
 - Vectors
 - Matrices
 - Dot Product
- Matrix Calculus for Deep Learning
 - Scalar derivatives
 - Vector derivatives
 - Matrix derivatives
 - DL context
- Probability Distributions for DL
 - Review of basic probability concepts
 - Key distributions

Why are Linear Algebra & Probability Critical for DL?

- Linear Algebra:
 - The "language" of neural networks.
 - Data represented as vectors and matrices.
 - Network operations (transformations, weights) are matrix multiplications.
- Probability:
 - Understanding data distributions.
 - Interpreting model outputs (e.g., classification probabilities).
 - Formulating loss functions and regularization.
 - Uncertainty modeling.




Linear Algebra Essentials

Vectors: Definition

- Ordered list of numbers (e.g., [x1, x2, x3]).
- In deep learning, vectors are commonly used to represent individual data points, features of an input, or learned representations (embeddings).
- Definition: A vector v of dimension n can be written as:

$$v = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix}$$

- This is a **column vector**
[conventionally used in DL for inputs and features]
 - A **row vector** would be $v^T = [v_1 \ v_2 \ \cdots \ v_n]$
- Geometric interpretation:
 - a point in n -dimensional space
 - or a directed line segment from the origin to that point.



Vectors: Operations

- Vector Addition: Element-wise sum of two vectors of the same dimension.

$$\begin{bmatrix} v_1 \\ v_2 \end{bmatrix} + \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} = \begin{bmatrix} v_1 + w_1 \\ v_2 + w_2 \end{bmatrix}$$

- Scalar Multiplication: Multiplying a vector by a scalar (a single number) scales its magnitude.

$$c \begin{bmatrix} v_1 \\ v_2 \end{bmatrix} = \begin{bmatrix} cv_1 \\ cv_2 \end{bmatrix}$$

Vectors: DL context

- An image can be "flattened" into a vector of pixel values
- A word can be represented by a "word embedding" vector
- The features describing a data point (e.g., height, weight, age) form a feature vector

Matrices

- Rectangular arrays of numbers (e.g., 2x3 matrix)
- Dimensions: `rows x columns`
- In deep learning, matrices are predominantly used to represent collections of data points (mini-batches), weights connecting layers in a neural network, or learned transformations.
- Definition: A matrix A with m rows and n columns (an $m \times n$ matrix) is written as:

$$A = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} \end{bmatrix}$$

Special Matrices

- **Identity matrix (I):** A square matrix with ones on the main diagonal and zeros elsewhere. When multiplied by another matrix, it leaves the matrix unchanged.

$$I = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

- **Diagonal Matrix:** A square matrix where all off-diagonal elements are zero.
- **Symmetric Matrix:** A square matrix where $A = A^T$ (transpose of A equals A).

Matrices: Operations

- **Matrix Addition/Subtraction:** Element-wise operations, requiring matrices to have the same dimensions.
- **Scalar Multiplication:** Multiplying every element of the matrix by a scalar.
- **Matrix Transpose (A^T):** Swapping rows and columns of a matrix. If A is $m \times n$, then A^T is $n \times m$.

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \implies A^T = \begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$$

Matrix Multiplication

- Not element-wise! Dot product of rows and columns
- **Definition:** The product of two matrices A (size: $m \times k$) and B (size: $k \times n$) results in a matrix C (size: $m \times n$). Each element c_{ij} of C is the dot product of the i -th row of A and the j -th column of B:

$$c_{ij} = \sum_{l=1}^k a_{il}b_{lj}$$

- **Rules for Multiplication:** For $A \times B$, the number of columns in A *must* equal the number of rows in B.

$$(m \times k) \times (k \times n) \rightarrow (m \times n)$$

Matrix Multiplication (contd.)

- **Non-Commutativity:** In general, $AB \neq BA$. The order of multiplication matters.
- **Geometric Interpretation:** Matrix multiplication can represent various geometric transformations such as scaling, rotation, reflection, and projection.
- **DL Context:**
 - The core operation within a neural network layer: $y=Wx+b$, where x is the input vector, W is the weight matrix, b is the bias vector, and y is the output vector.
 - Processing mini-batches: If X is a matrix where each row is a data point (or each column, depending on convention), and W is a weight matrix, then XW or WX processes the entire batch efficiently.

Dot Product (Vector Inner Product)

- **Definition:** For two vectors $v = \begin{bmatrix} v_1 \\ \vdots \\ v_n \end{bmatrix}$ and $w = \begin{bmatrix} w_1 \\ \vdots \\ w_n \end{bmatrix}$ of the same dimension n , their dot product is a scalar:

$$v \cdot w = v^T w = \sum_{i=1}^n v_i w_i = v_1 w_1 + v_2 w_2 + \cdots + v_n w_n$$

- **Geometric Interpretation:**
 - If $v \cdot w = 0$, the vectors are orthogonal (perpendicular).
 - It relates to the angle θ between vectors: $v \cdot w = \|v\| \cdot \|w\| \cos(\theta)$. This is used for cosine similarity, a measure of how similar two vectors are.

Dot Product (Vector Inner Product) (contd.)

- DL Context:
 - The "weighted sum" in a neuron's activation: The dot product of input features x and weights w (i.e., $w \cdot x$) forms the core of many activation functions before a non-linearity is applied.
 - Attention mechanisms in advanced architectures use dot products to compute similarity scores between different parts of the input.



Matrix Calculus for Deep Learning

Matrix Calculus - Introduction

- Why is this hard but crucial?
 - Deep Learning involves optimizing functions (loss functions) with millions of parameters.
 - We need to compute gradients to update these parameters efficiently.
 - "Backpropagation" is essentially repeated application of the chain rule.
- Review:
 - Scalar Derivatives: $\frac{dy}{dx}$ (slope)
 - Chain Rule: $\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx}$ (how changes propagate)

Matrix Calculus - Vector Derivatives

- **Gradient ($\nabla_x f(x)$):** Derivative of a **scalar** function (f) with respect to a **vector** input (x).
- Result: A vector of partial derivatives.

- Let $f(x)$ be a scalar function where $x = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$. The gradient of f with respect to x is a vector containing all the partial derivatives of f with respect to each component of x :

$$\nabla_x f(x) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$

Vector Derivatives (contd.)

- **Interpretation:** The gradient vector points in the direction of the steepest increase of the function f . In optimization algorithms like gradient descent, we move in the opposite direction of the gradient to find a minimum.
- **DL Context:** This is used to compute the gradients of the *loss function* (a scalar value) with respect to the network's *parameters* (weights and biases, which are vectors or matrices). For instance, $\nabla_w \mathcal{L}$ tells us how to adjust the weight vector w to reduce the loss \mathcal{L} .

Vector Derivatives (contd.)

- **Jacobian Matrix (\mathbf{J}):** Derivative of a vector-valued function with respect to a vector input.

- Let $y = f(x)$ be a function that maps an n -dimensional input vector x to an m -

dimensional output vector $y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$. The Jacobian matrix of f with respect to x is an

$m \times n$ matrix:

$$\mathbf{J} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_1}{\partial x_2} & \cdots & \frac{\partial y_1}{\partial x_n} \\ \frac{\partial y_2}{\partial x_1} & \frac{\partial y_2}{\partial x_2} & \cdots & \frac{\partial y_2}{\partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y_m}{\partial x_1} & \frac{\partial y_m}{\partial x_2} & \cdots & \frac{\partial y_m}{\partial x_n} \end{bmatrix}$$

Vector Derivatives (contd.)

- **DL Context:** Jacobians are essential for understanding how errors propagate between layers. If we have a sequence of operations like $z = g(y)$ and $y = f(x)$, then $\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx}$ (using the chain rule), where these "derivatives" are actually Jacobian matrices.

Matrix Derivatives

- In deep learning, we frequently need to compute the derivative of a scalar loss function with respect to a matrix of weights.
- **Derivative of a scalar with respect to a matrix:**

Let $f(X)$ be a scalar function where X is an $m \times n$ matrix. The derivative of f with respect to X is an $m \times n$ matrix where each element (i, j) is $\frac{\partial f}{\partial X_{ij}}$:

$$\frac{\partial f}{\partial X} = \begin{bmatrix} \frac{\partial f}{\partial X_{11}} & \frac{\partial f}{\partial X_{12}} & \cdots & \frac{\partial f}{\partial X_{1n}} \\ \frac{\partial f}{\partial X_{21}} & \frac{\partial f}{\partial X_{22}} & \cdots & \frac{\partial f}{\partial X_{2n}} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f}{\partial X_{m1}} & \frac{\partial f}{\partial X_{m2}} & \cdots & \frac{\partial f}{\partial X_{mn}} \end{bmatrix}$$


Backpropagation as Chain Rule Application

Steps

epoch
 ↑
 loop
 100 ← 1000
 1) for i in range(q):
 a) 1 student \rightarrow FP \rightarrow
 b) Loss (mse)
 c) Adjust all w, b

$$w_{\text{new}} = w_{\text{old}} - \eta \frac{\partial L}{\partial w_{\text{old}}}$$

 q times





Probability Distributions for Deep Learning

Probability Review - Basics

- **Random Variables (RV):** Outcomes of random phenomena.
 - Discrete: Countable outcomes (e.g., coin flip).
 - Continuous: Infinite outcomes within a range (e.g., height, temperature).
- **Probability Mass Function (PMF):** For discrete RVs, it gives the probability that the random variable takes on a specific value. $P(X = x)$.
- **Probability Density Function (PDF):** For continuous RVs, it describes the likelihood of the random variable falling within a particular range of values. The probability of X being in an interval $[a, b]$ is $\int_a^b f(x)dx$.

Probability Review – Basics (contd.)

- **Expectation ($E[X]$):** The average or mean value of a random variable.
 - For discrete RV: $E[X] = \sum xP(X=x)$.
 - For continuous RV: $E[X] = \int xf(x)dx$.
- **Variance ($Var[X]$):** A measure of how spread out the values of a random variable are from its mean.
 - $Var[X] = E[(X - E[X])^2]$.
- Standard Deviation is $\sigma = \sqrt{Var[X]}$

Key Probability Distributions for DL - Bernoulli

- Describes a single experiment with only two possible outcomes (e.g., success/failure, 0/1), where the probability of "success" is p .
- **PMF:**

$$f(x) = \begin{cases} p, & \text{if } k = 1 \\ 1 - p, & \text{if } k = 0 \end{cases}$$

This can also be written as $P(X = k) = p^k(1 - p)^{1-k}$

- DL context:
 - **Binary Classification:** Output of a sigmoid layer can be interpreted as p .
 - **Dropout:** Each neuron is independently "dropped" (set to 0) with a Bernoulli probability.

Key Probability Distributions for DL - Categorical

- A generalization of the Bernoulli distribution for discrete random variables with more than two possible outcomes (e.g., rolling a die, classifying an image into one of several categories).
- It has k mutually exclusive possible outcomes, each with a specific probability p_i , where $\sum_{i=1}^k p_i = 1$
- **PMF:** $P(X = i) = p_i$ for $i \in \{1, 2, \dots, k\}$.
- DL context:
 - **Multi-class Classification:** Output of a softmax layer gives probabilities for each class.
 - **Language Modeling:** Predicting the next word from a vocabulary.

Key Probability Distributions for DL - Normal (Gaussian)

- The "bell curve" distribution.
- Parameters: Mean (μ) and Standard Deviation (σ).

- PDF:

$$f(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

- DL context:
 - **Weight Initialization:** Often initialized from Normal distributions (e.g., Xavier, He).
 - **Generative Models:** VAEs often model latent spaces or outputs as Gaussian.
 - **Loss Functions:** L2 Loss (MSE) is related to maximizing likelihood under Gaussian noise.
 - **Regularization:** Adding Gaussian noise to inputs/activations.

Summary

- Linear Algebra:
 - **Matrix Multiplication:** The fundamental operation of neural networks.
 - **Matrix Calculus / Gradients:** How neural networks learn (backpropagation).
 - **Chain Rule:** The mathematical backbone of backpropagation.
- Probability:
 - **Bernoulli:** Binary outcomes, dropout.
 - **Categorical (Softmax):** Multi-class classification.
 - **Normal:** Weight initialization, generative models, noise.