



Topic 6a

Convolutional Neural Network (Basics)

CSE465: Pattern Recognition and Neural Network

Sec: 3

Faculty: Silvia Ahmed (SvA)

Summer 2025

Topics

1. What is a Convolutional Neural Network (CNN)?
2. Basic intuition.
3. Visual Cortex vs CNN
4. Operations:
 1. Convolution
 2. Padding
 3. Stride
 4. Pooling
5. CNN Architecture

What is CNN?

- Convolutional Neural Network, also known as convnet, or CNNs, are a special kind of neural network for processing data that has a known grid-like topology like time series data (1D) or images (2D).

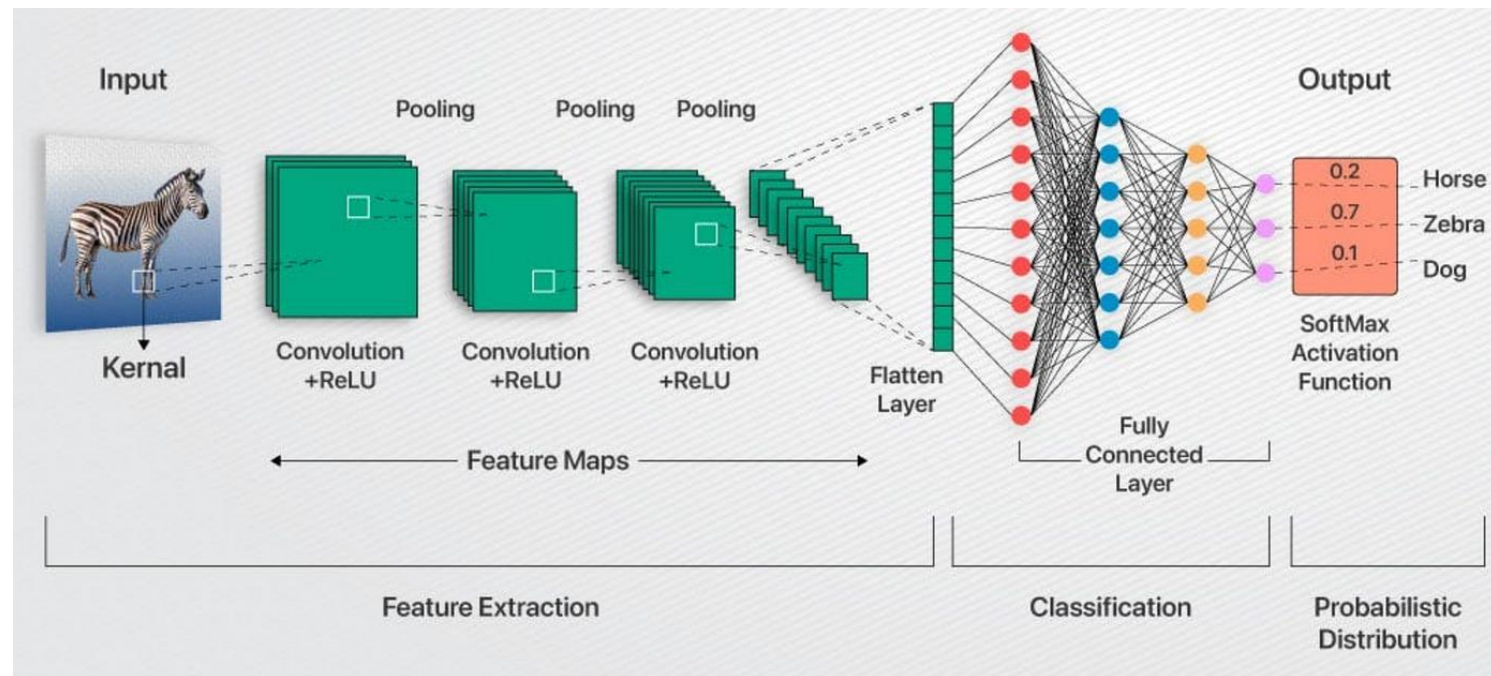
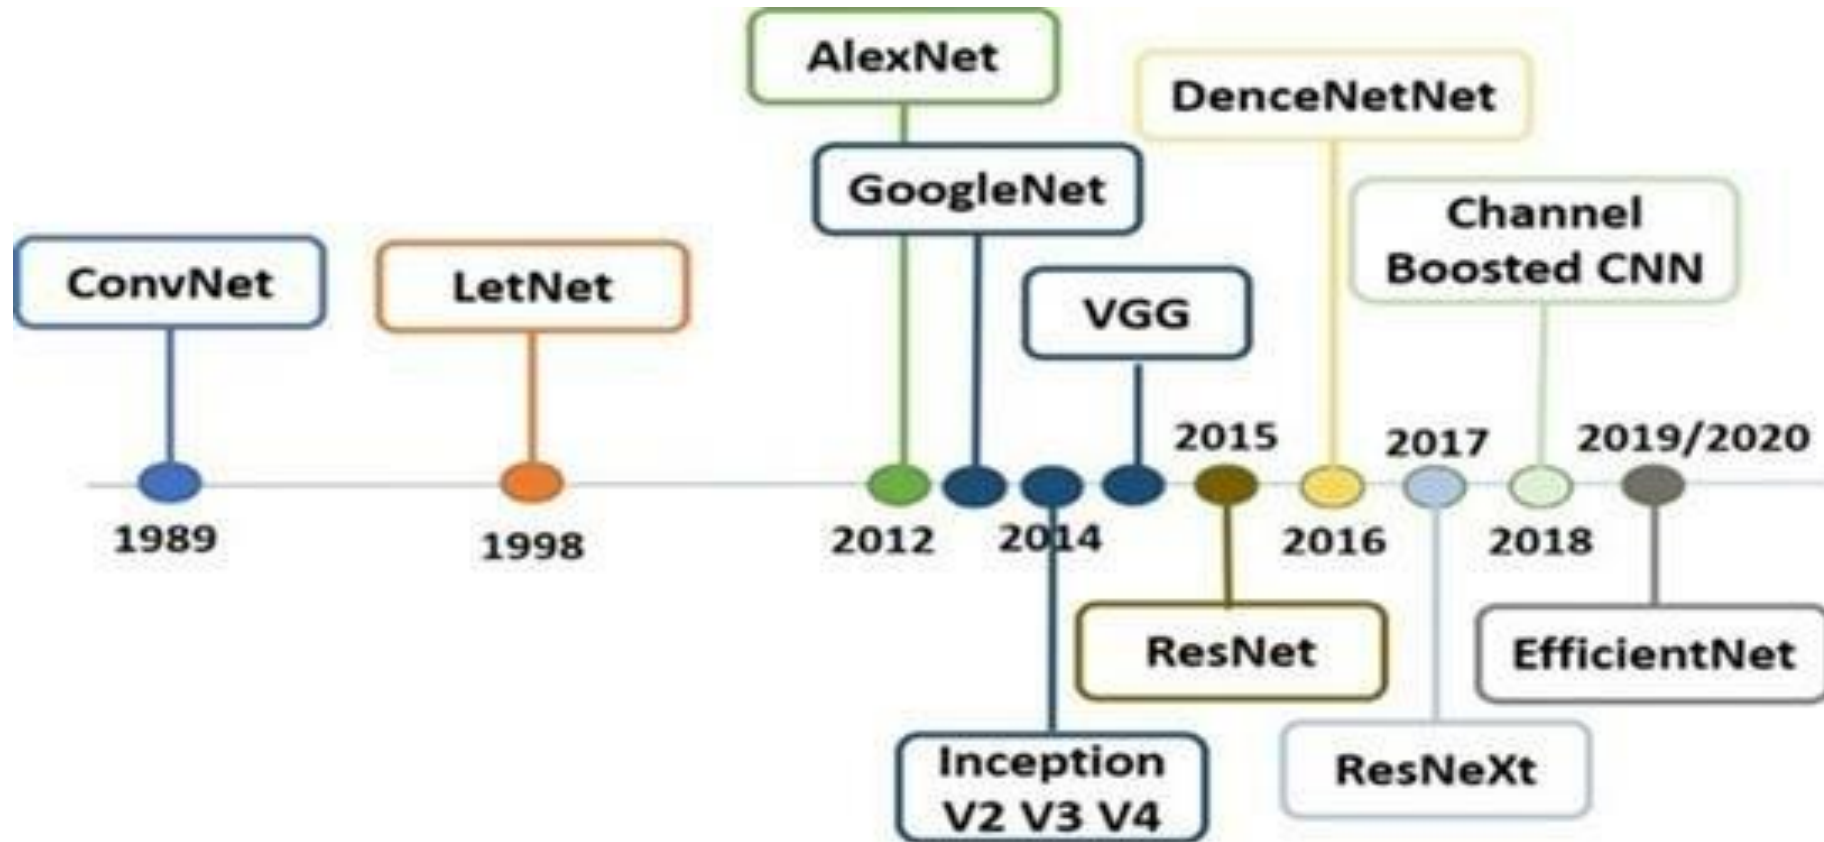


Figure: Basic structure of a CNN [1]

Layers in CNN


- A special layer for “convolution” operation.
- If there’s even a single “convolution” layer in a neural network, then that becomes a CNN.
- 3 types of layers:
 - Convolution layer
 - Pooling layer
 - Fully-connected (FC) layer

Brief history




Limitations of ANN

- High computational cost
- Overfitting
- Loss of important info (e.g., spatial arrangement of pixels)





- 28x28 pixels grayscale MNIST image
- To feed into ANN, this must be flattened out to make 784 input nodes
- So layer 1 with 100 neurons will require learning parameters of 78400, and so on.
- Images with higher dimension will make this number exponentially bigger.
- Learning this huge no of parameters will take a lot of time which will eventually end up in overfitting.
- Flattening of the 2D structure loses spatial information as well.

Intuition





Visual Cortex



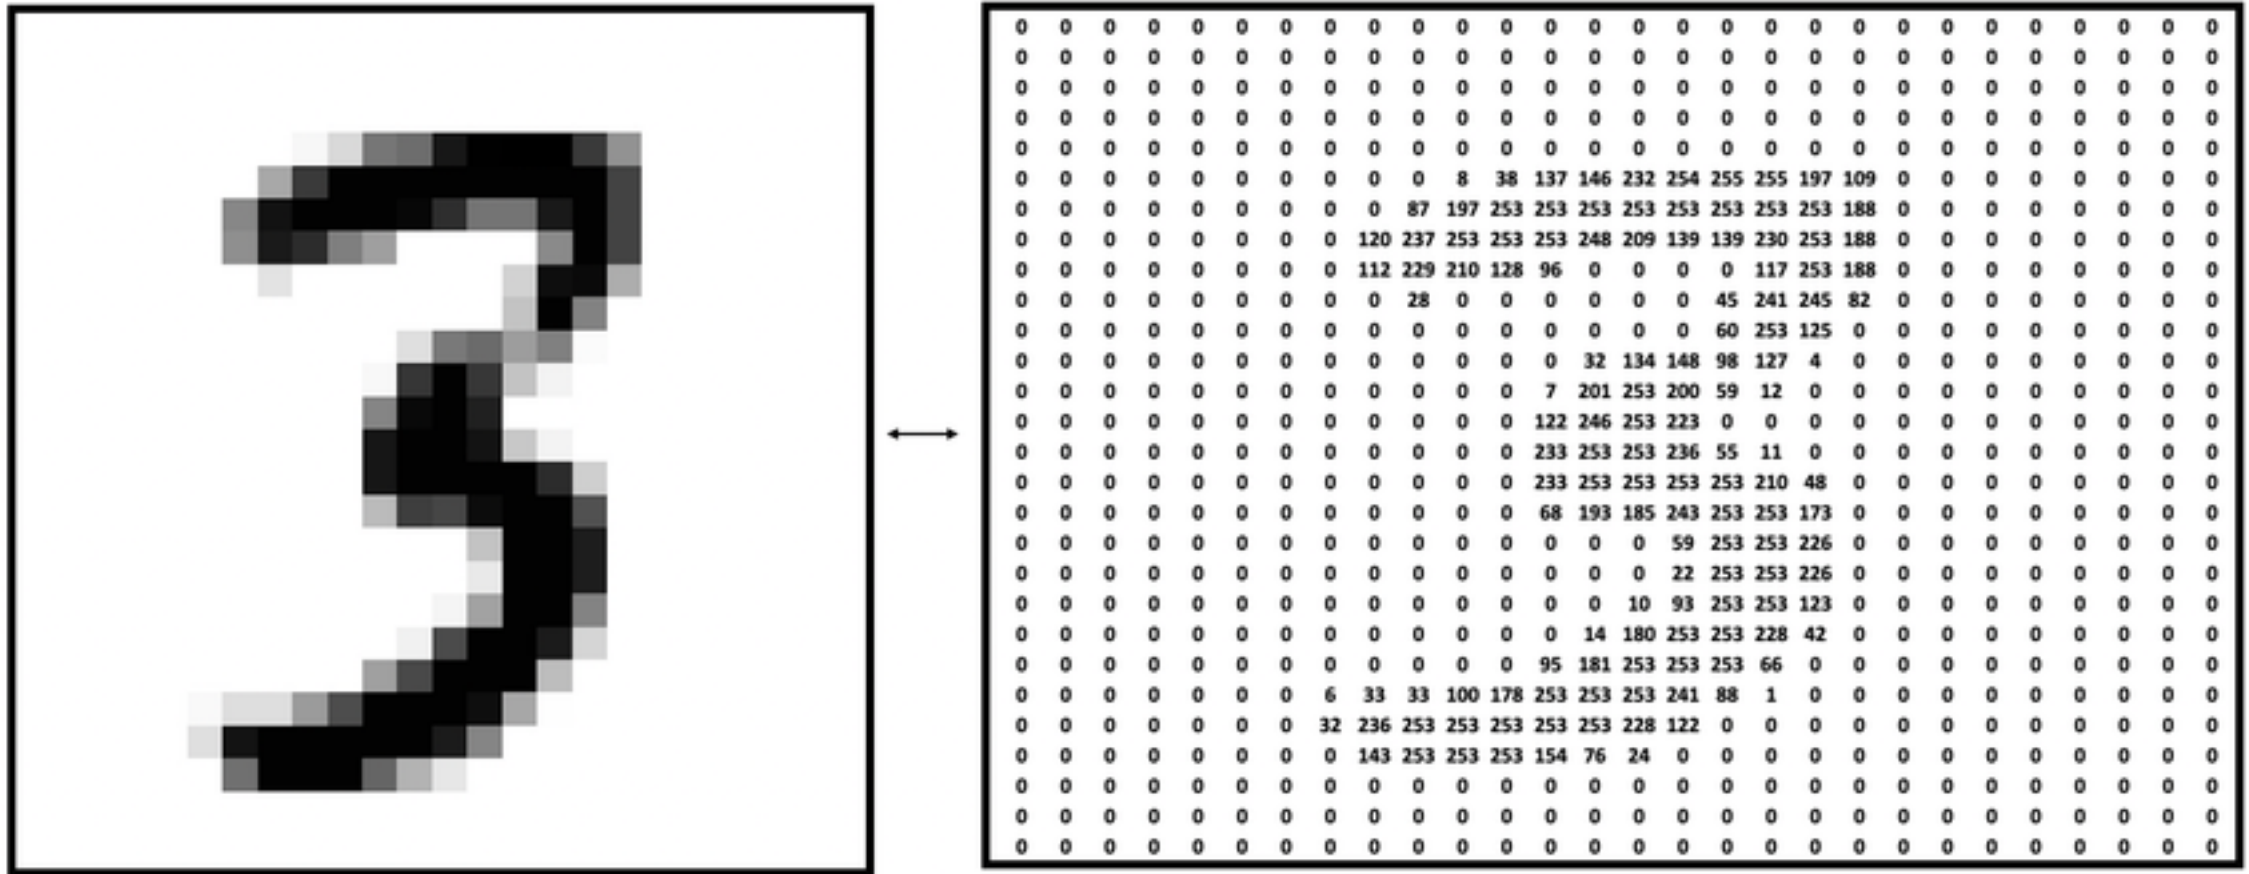
Hubel and Wiesel Redux

A Experimental setup

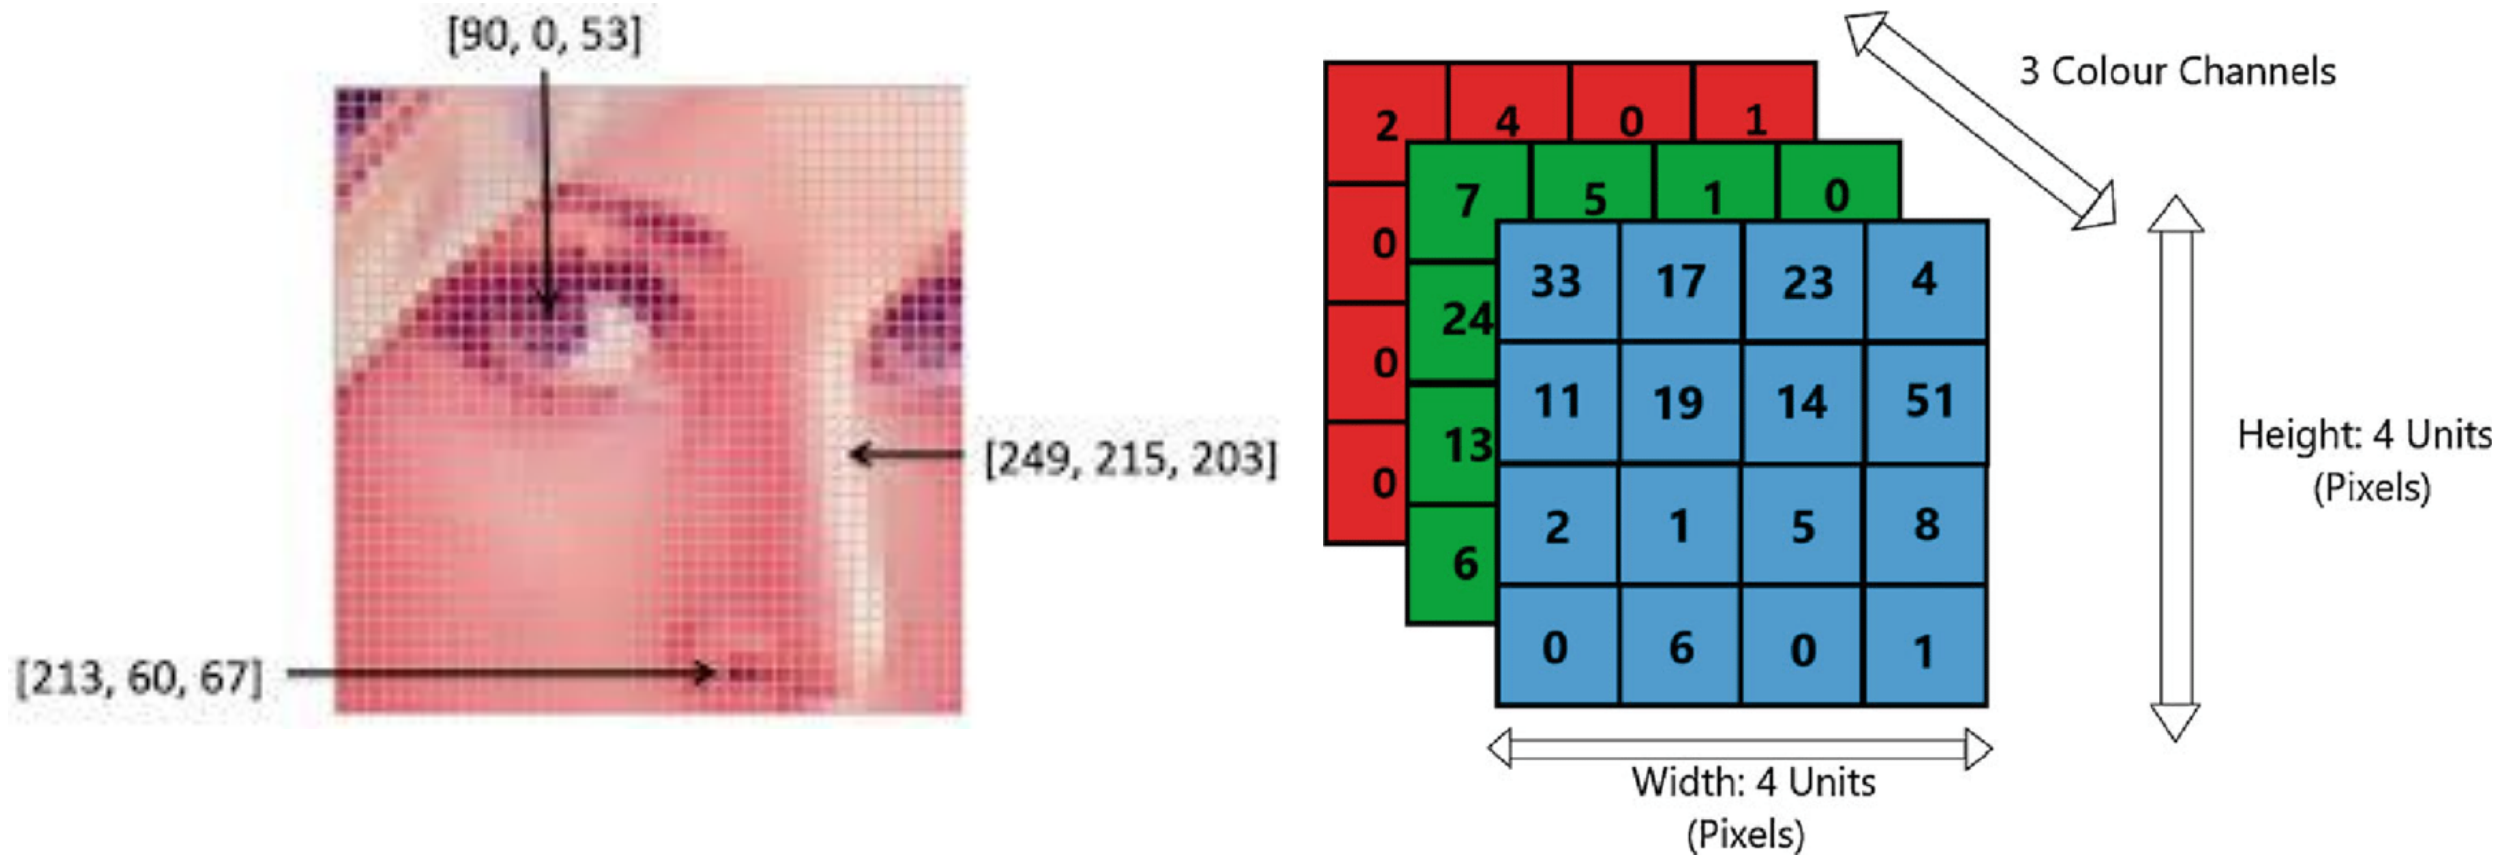


- Their original findings, showing that neurons in V1 detect simple edge-like patterns, while later layers respond to increasingly complex features, have been largely validated by modern neuroscience.
- Hubel and Wiesel categorized neurons into **simple** (edge detectors) and **complex** (more spatially invariant feature detectors).

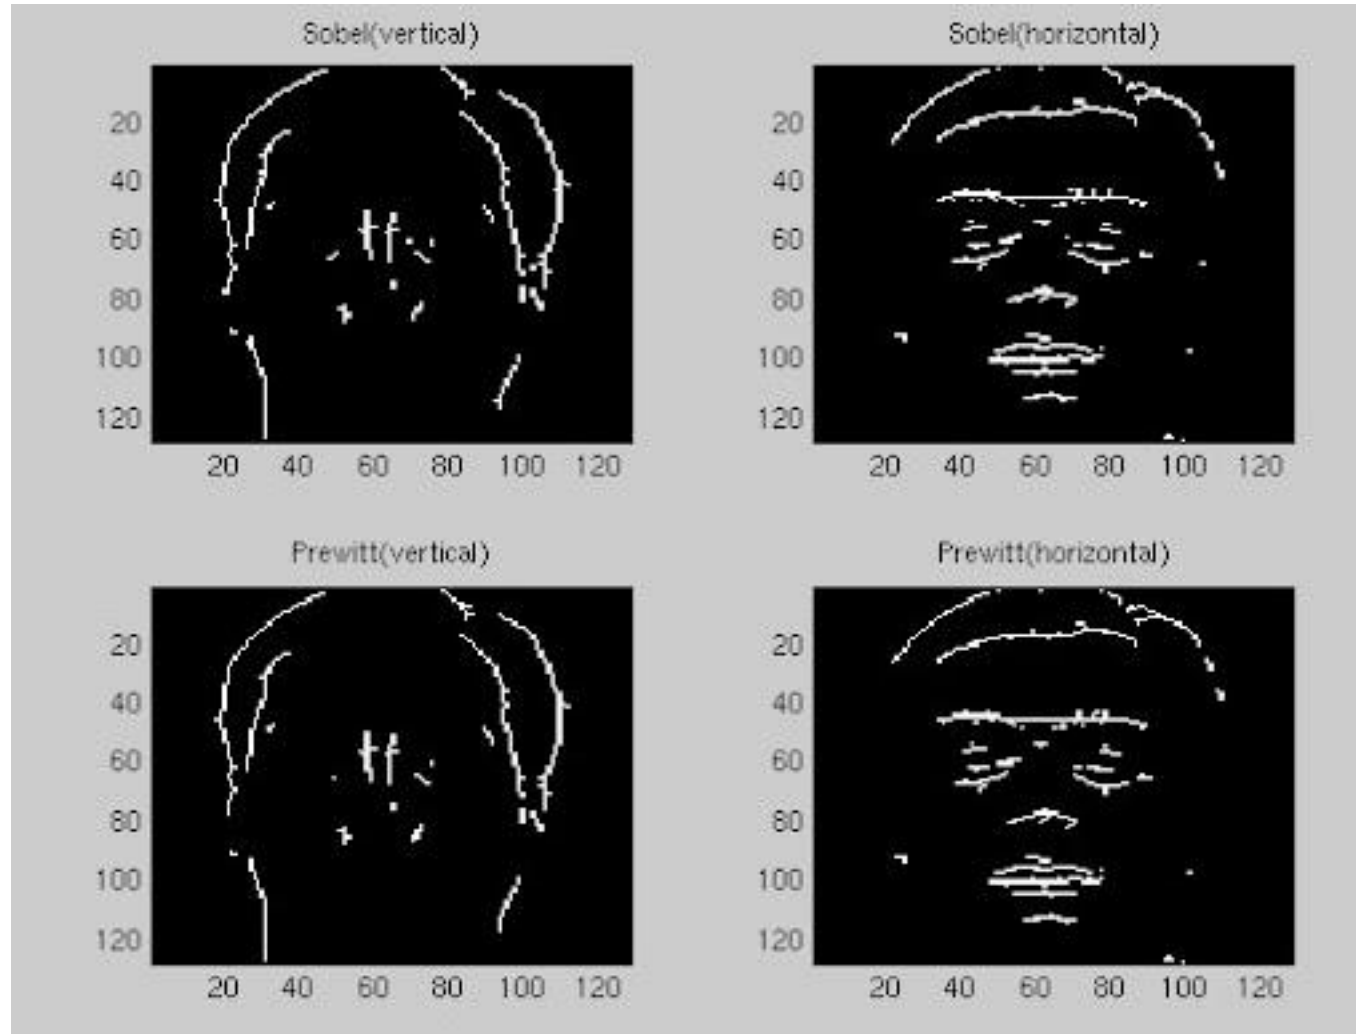
Basics of Images (Grayscale images)



Basics of Images (Color images)



Edge Detection (Convolution operation)



Edges are changed intensity

Horizontal Edge Detection

0	0	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0
255	255	255	255	255	255
255	255	255	255	255	255
255	255	255	255	255	255



Image

*

-1	-1	-1
0	0	0
1	1	1




Filter/Kernel


=

Feature Map


Horizontal Edge Detection (calculation)




Horizontal Edge Detection (calculation)




Horizontal Edge Detection (calculation)




Horizontal Edge Detection (calculation)




Horizontal Edge Detection (calculation)




Horizontal Edge Detection (calculation)




Horizontal Edge Detection (calculation)




Horizontal Edge Detection (calculation)



Horizontal Edge Detection (calculation)



Horizontal Edge Detection (calculation)



Live Demonstration

- <https://deeplizard.com/resource/pavq7noze2>
- Red: Positive activation
- Blue: Opposite edge detection
- For example, if we chose left-edge filter, then red means left edges are detected. On the other hand, blue means opposite (right-edge) has been detected.

Size of feature map

0	0	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0
255	255	255	255	255	255
255	255	255	255	255	255
255	255	255	255	255	255

$n \times n$

*

-1	-1	-1
0	0	0
1	1	1


$k \times k$

=

0	0	0	0
255	255	255	255
255	255	255	255
0	0	0	0

$(n - k + 1) \times (n - k + 1)$

Working with RGB images



Issues with Convolution

- Reduction in Spatial Dimensions:
 - the output size of the feature map decreases after each convolutional layer due to the kernel sliding over the input image.
 - If we use a $k \times k$ filter on an $n \times n$ input with a stride of 1, the output size is reduced to $(n - k + 1) \times (n - k + 1)$.
 - This means that as we go deeper into the network, the feature maps keep shrinking, leading to information loss.
 - Example: For a 28×28 input with a 3×3 filter and no padding:
Output size = $28 - 3 + 1 = 26$
 - Each layer further reduces the size, which can lead to vanishing spatial information in deep networks.

Issues with Convolution (contd.)

- Loss of Edge Information:
 - **edge pixels** are used fewer times compared to central pixels during convolution, leading to **biased feature extraction**.
 - Padding ensures that even edge and corner features contribute equally in the learning process.
- No Control Over Output Size
- **Solution: Add extra pixels** (usually zeros) around the input to maintain or control the spatial dimensions.

Padding

- Padding refers to **adding extra pixels** (usually zeros) around the input image before applying convolution operations.
- It helps control the spatial size of feature maps and enhances model performance.
- Benefits:
 - Maintains Spatial Dimensions
 - Prevents Loss of Edge Information
 - Allows for Control Over Feature Map Size
 - Improves Performance in Deep CNNs

Types of Padding in CNNs

1. Valid Padding ("No Padding"):

- **Definition:** No extra pixels are added, meaning the kernel applies only to the original image.
- **Effect:** The feature map shrinks after each convolution.
- **Formula:**

$$\text{Output size} = (n - k + 1) \times (n - k + 1)$$

- **Example:**

- **Input:** 28 x 28, **Filter:** 3 x 3, No Padding
 - **Output:** $(28 - 3 + 1) \times (28 - 3 + 1) = 26 \times 26$
- When to use?
 - When reducing spatial size is acceptable (e.g., classification tasks).
 - When deeper layers apply **global average pooling** (e.g., ResNet, MobileNet).

Types of Padding in CNNs (contd.)

2. Same Padding (Zero-Padding):

- **Definition:** Padding is added to ensure that the **output size is the same as the input size**.
- **Formula for padding size:**

$$P = \frac{(k-1)}{2} \text{ (for odd-sized kernels)}$$

- **Example:**
 - **Input:** 28 x 28, **Filter:** 3 x 3, **Padding:** $\frac{3-1}{2} = 1$ pixel
 - **Output:** 28 x 28 (unchanged)
- **Advantages:**
 - Keeps feature map size constant, simplifying architecture design.
 - Useful for deep networks like **VGG**, **ResNet**.

Types of Padding in CNNs (contd.)

3. Full Padding:

- **Definition:** Maximum padding is applied so that every pixel gets covered by the kernel the same number of times.
- **Formula for padding size, $P = k - 1$**
- **Formula for Output Size:**
$$\text{Output size} = (n + 2P - k + 1) \times (n + 2P - k + 1)$$
- **Example:**
 - **Input:** 28 x 28, **Filter:** 3 x 3, **Full padding:** 2 pixels
 - **Output:** $(28 + 2 \times 2 - 3 + 1)$ or 30 x 30
- **When to use?**
 - If we want **larger feature maps** than the input size.
 - Used in **certain styles of generative models (GANs, autoencoders).**

Types of Padding Techniques

1. Zero Padding (Most Common)

- Adds zeros around the image.
- Simple and widely used.

0	0	0	0	0	0
0	1	2	3	4	0
0	5	6	7	8	0
0	9	10	11	12	0
0	13	14	15	16	0
0	0	0	0	0	0

2. Replication Padding

- Duplicates edge values to preserve texture.
- Used in **image super-resolution**.

1	1	2	3	4	4
1	1	2	3	4	4
5	5	6	7	8	8
9	9	10	11	12	12
13	13	14	15	16	16
13	13	14	15	16	16

Types of Padding Techniques

3. Reflection Padding

- Mirrors pixels at the edge..
- Reduces border artifacts in **image processing**.

6	5	6	7	8	7
2	1	2	3	4	3
6	5	6	7	8	7
10	9	10	11	12	11
14	13	14	15	16	15
10	9	10	11	12	11

2. Circular Padding

- Wraps the image around itself.
- Used in **periodic signal processing**.

16	13	14	15	16	13
4	1	2	3	4	1
8	5	6	7	8	5
12	9	10	11	12	9
16	13	14	15	16	13
4	1	2	3	4	1

Stride

- **Stride** is the step size by which the convolutional filter (kernel) moves across the input image during convolution.
- It determines:
 - How much the receptive field moves at each step
 - How much the output shrinks compared to the input
 - How much computational efficiency is improved
- **$S = 1$** → The filter moves **one pixel** at a time → **Dense feature extraction**
- **$S = 2$** → The filter moves **two pixels** at a time → **Downsampling occurs**
- **$S > 2$** → The filter moves **more than two pixels** at a time → **Aggressive Downsampling occurs**

Formulation of Padding and Strides

- For a **stride S** and **padding P**, the **output size** of a convolutional layer is given by:

$$\text{Output width} = \left\lfloor \frac{(\text{Input width} + 2P - k)}{S} \right\rfloor + 1$$

$$\text{Output height} = \left\lfloor \frac{(\text{Input height} + 2P - k)}{S} \right\rfloor + 1$$

Where:

- k = kernel/filter size
- S = Stride
- P = Padding
- Input width/height = Original Image size

Further Issues with Convolution

1. Memory issues: Example, For a **224 × 224 RGB image**, assuming:

- **64 feature maps** in a layer,
- **32-bit float representation (4 bytes per value),**

the memory required for one layer is:

$$224 \times 224 \times 64 \times 4 \text{ bytes} = 12.8 \text{ MB}$$

- For **10 layers**, the memory usage becomes **128 MB per image!**
- With **batch processing**, memory usage increases further

2. Translation variance:

- CNNs are **not fully translation-invariant**, meaning:
 - If an object shifts slightly in an image, the CNN might classify it differently.
 - Small shifts cause different activations, affecting feature maps.
 - Example: Digit Recognition (MNIST Dataset): If a "5" is shifted **one pixel to the right**, the CNN may **misclassify it as "3"** due to different activations.



Solution:
Pooling

Pooling (layer)

- Pooling is a **downsampling operation** used in CNNs to reduce the spatial dimensions of feature maps while preserving important information.
- It helps in:
 - **Reducing computational cost** by shrinking the feature map size
 - **Improving translation invariance** (i.e., detecting patterns regardless of their exact location)
 - **Preventing overfitting** by reducing unnecessary details
- How Pooling Works?
 - Pooling operates on small regions (typically **2×2**) of the feature map and summarizes them using a specific function

Types of Pooling

1. Max Pooling (Most Common)

- Takes the **largest value** in the pooling window.
- Preserves the **strongest features** (e.g., edges, textures).
- Commonly used in deep CNN architectures (e.g., VGG, ResNet).

1	3	2	4
5	6	8	7
9	10	12	11
13	14	16	15



2 x 2 max pooling,
Stride = 1


6	8
14	16

- Max pooling reduces noise and keeps dominant features
- Commonly used after convolutional layers to reduce dimensions

Types of Pooling (contd.)

2. Average Pooling:

- Computes the **mean value** in the pooling region.
- Retains **global structure** but loses some sharp details.
- Used in **shallow networks or specific tasks like regression**.



- **Blurs sharp edges** but **keeps overall distribution**
- Used in classification tasks like ImageNet models (AlexNet, ResNet)
- Better for smooth feature extraction

Types of Pooling (contd.)

3. Global Pooling (Global Average Pooling):

- Reduces the entire feature map to a **single value per channel**.
- Used in architectures like **Google's Inception** and **ResNet**.
- Helps replace **fully connected (FC) layers**, reducing parameters.

1	3	2	4
5	6	8	7
9	10	12	11
13	14	16	15




8.5
(Single value per channel)

- Used before the final classification layer in **ResNet**
- Eliminates need for large FC layers, reducing parameters
- Prevents overfitting in deep networks

Benefits of Pooling

- Reduced size:



- Translation invariance:





Figure 14-9. Invariance to small translations

Benefits of Pooling (contd.)

- Enhanced Features: Benefits of Pooling
 - Only in case of Max pooling
- No need of training



Stride & Pooling

- Pooling layers typically use **stride = pool size** to ensure:
 - **Non-overlapping** receptive fields (e.g., **2×2 pool, stride = 2**)
 - Downsampling without overlap
 - Smaller, more efficient feature maps

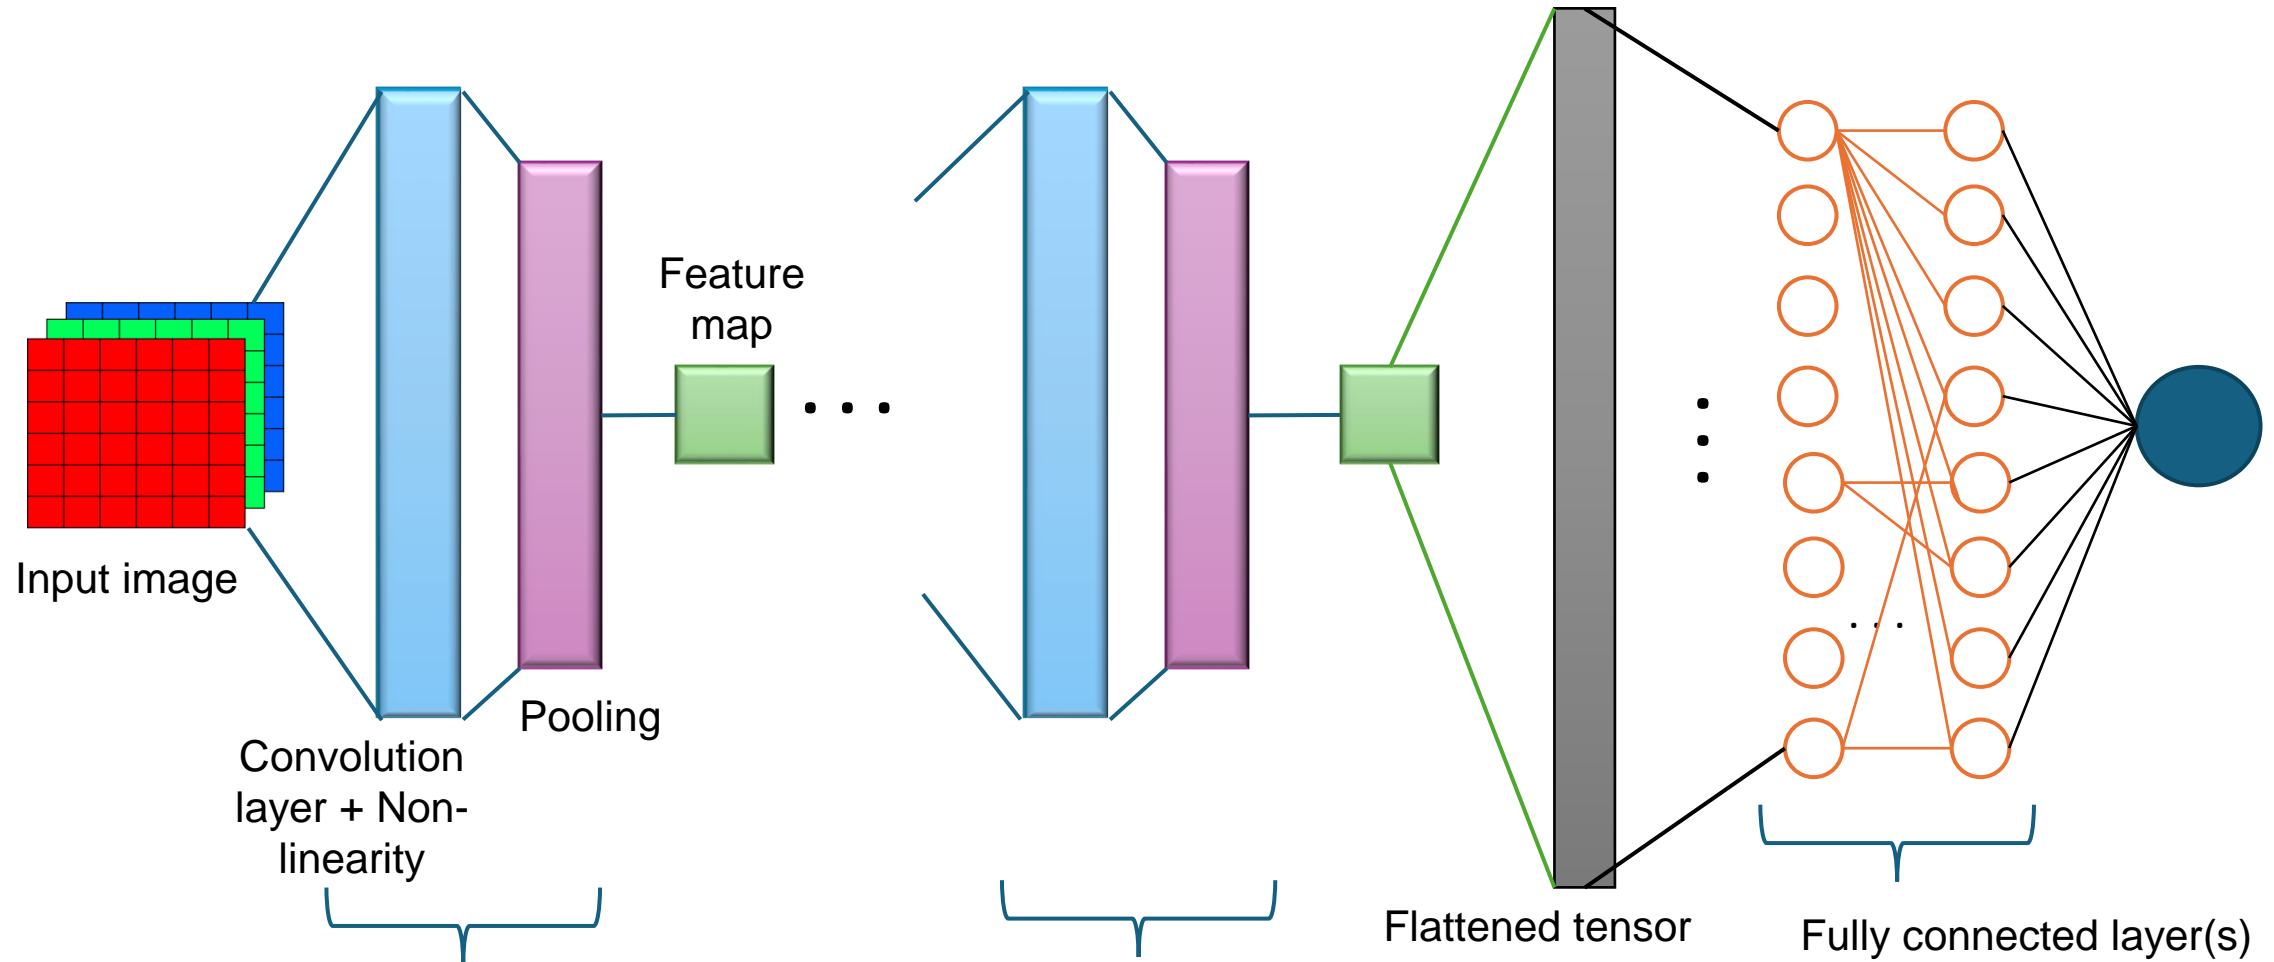
Pooling vs Convolution: Key Differences

Feature	Convolution	Pooling
Purpose	Extracts features (edges, textures)	Reduces feature map size
Operation	Learns from data (weights)	Fixed function (max/avg)
Effect	Preserves information	Removes redundant information
Computational Cost	High	Low

When NOT to Use Pooling


- **If spatial relationships are important** (e.g., segmentation tasks)
- **If information loss is harmful** (e.g., GANs, super-resolution models)
- **If using Strided Convolution as an alternative** (e.g., ResNet, MobileNet)

Basic CNN architecture




Difference in CNN Architecture

- Number of Convolution layer
- Number of filters/kernels
- Stride
- Pooling
- Number of Fully Connected (FC) nodes
- Number of FC layers
- Activation functions
- Dropouts
- Batch norm



Example: LeNet-5



Source: [3]

References

- [1] <https://ravjot03.medium.com/decoding-cnns-a-beginners-guide-to-convolutional-neural-networks-and-their-applications-1a8806cbf536>
- [2] “Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow”, 2nd Edition
- [3] <https://www.analyticsvidhya.com/blog/2021/03/the-architecture-of-lenet-5/>
- [4] Youtube playlist: 100 Days of Deep Learning by CampusX